

es21btech11025-assign6

March 14, 2024

EP4130: Data Science Analysis

Assignment 6

Ranveer Sahu- ES21BTECH11025

```
[1]: import numpy as np
import scipy as sp
import matplotlib.pyplot as plt
import emcee
!pip install astroml
from astroML.plotting import plot_mcmc
```

Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: astroml in
/home/coder/.local/lib/python3.10/site-packages (1.0.2.post1)
Requirement already satisfied: scikit-learn>=0.18 in
/home/coder/.local/lib/python3.10/site-packages (from astroml) (1.4.1.post1)
Requirement already satisfied: numpy>=1.13 in
/home/coder/.local/lib/python3.10/site-packages (from astroml) (1.24.3)
Requirement already satisfied: matplotlib>=3.0 in
/home/coder/.local/lib/python3.10/site-packages (from astroml) (3.7.1)
Requirement already satisfied: scipy>=0.18 in
/home/coder/.local/lib/python3.10/site-packages (from astroml) (1.10.1)
Requirement already satisfied: astropy>=3.0 in
/home/coder/.local/lib/python3.10/site-packages (from astroml) (6.0.0)
Requirement already satisfied: astropy-iers-data>=0.2023.10.30.0.29.53 in
/home/coder/.local/lib/python3.10/site-packages (from astropy>=3.0->astroml)
(0.2024.3.11.18.33.3)
Requirement already satisfied: packaging>=19.0 in
/home/coder/.local/lib/python3.10/site-packages (from astropy>=3.0->astroml)
(23.0)
Requirement already satisfied: pyerfa>=2.0 in
/home/coder/.local/lib/python3.10/site-packages (from astropy>=3.0->astroml)
(2.0.1.1)
Requirement already satisfied: PyYAML>=3.13 in /usr/lib/python3/dist-packages
(from astropy>=3.0->astroml) (5.4.1)

```

Requirement already satisfied: contourpy>=1.0.1 in
/home/coder/.local/lib/python3.10/site-packages (from matplotlib>=3.0->astroml)
(1.0.7)
Requirement already satisfied: fonttools>=4.22.0 in
/home/coder/.local/lib/python3.10/site-packages (from matplotlib>=3.0->astroml)
(4.39.3)
Requirement already satisfied: python-dateutil>=2.7 in
/home/coder/.local/lib/python3.10/site-packages (from matplotlib>=3.0->astroml)
(2.8.2)
Requirement already satisfied: pillow>=6.2.0 in /usr/lib/python3/dist-packages
(from matplotlib>=3.0->astroml) (9.0.1)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/lib/python3/dist-
packages (from matplotlib>=3.0->astroml) (2.4.7)
Requirement already satisfied: cycycler>=0.10 in
/home/coder/.local/lib/python3.10/site-packages (from matplotlib>=3.0->astroml)
(0.11.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
/home/coder/.local/lib/python3.10/site-packages (from matplotlib>=3.0->astroml)
(1.4.4)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/home/coder/.local/lib/python3.10/site-packages (from scikit-
learn>=0.18->astroml) (3.3.0)
Requirement already satisfied: joblib>=1.2.0 in
/home/coder/.local/lib/python3.10/site-packages (from scikit-
learn>=0.18->astroml) (1.3.2)
Requirement already satisfied: six>=1.5 in /usr/lib/python3/dist-packages (from
python-dateutil>=2.7->matplotlib>=3.0->astroml) (1.16.0)

/home/coder/.local/lib/python3.10/site-
packages/astroML/linear_model/linear_regression_errors.py:10: UserWarning:
LinearRegressionwithErrors requires PyMC3 to be installed
  warnings.warn('LinearRegressionwithErrors requires PyMC3 to be installed')

```

Question 1

In 1919, two expeditions sailed from Britain to test if the light deflection from stars agrees with Einstein's General Theory of Relativity. Einstein's theory predicts a value of 1.74 arc-seconds, whereas Newtonian gravity predicts a value exactly half of that. The team by Eddington obtained a value of 1.61 ± 0.40 arc-seconds, while the team by Crommelin reported a value of 1.98 ± 0.16 arc-seconds. Calculate the Bayes factor between General Relativity and Newtonian gravity from those data, assuming Gaussian likelihoods.

given data's:

```

[2]: einstein_time = 1.74
     newton_time = einstein_time / 2
     eddington_time = 1.61
     eddington_error = 0.40

```

```
crommelin_time = 1.98
crommelin_error = 0.16
```

```
[3]: einstein_evidence = sp.stats.norm(einstein_time, eddington_error).
      ↪pdf(eddington_time) * \
      sp.stats.norm(einstein_time, crommelin_error).
      ↪pdf(crommelin_time)
```

```
[4]: newton_evidence = sp.stats.norm(newton_time, eddington_error).
      ↪pdf(eddington_time) * \
      sp.stats.norm(newton_time, crommelin_error).
      ↪pdf(crommelin_time)
```

```
[5]: print("Bayes Factor for the models of Einstein and Netwon are {} in favour of_
      ↪Einstein's model.".format(einstein_evidence/newton_evidence))
```

Bayes Factor for the models of Einstein and Netwon are 48164622958.34179 in favour of Einstein's model.

Question 2

For exercise 1 in arXiv:1008.4686, calculate the 68% and 95% joint confidence intervals on b and m .

```
[6]: data = [[203,495,21,],
             [58,173,15,],
             [210,479,27,],
             [202,504,14,],
             [198,510,30,],
             [158,416,16,],
             [165,393,14,],
             [201,442,25,],
             [157,317,52,],
             [131,311,16,],
             [166,400,34,],
             [160,337,31,],
             [186,423,42,],
             [125,334,26,],
             [218,533,16,],
             [146,344,22]]

data = np.array(data)

x = data[:, 0]
y = data[:, 1]
e = data[:, 2]
```

As we learn in Machine learning course: models functions given below

1. linear model

```
[7]: def linear_model(x , theta):  
      return theta[0] + theta[1] * x
```

2. log prior model

```
[8]: def log_prior(theta):  
  
      if theta[2] < 0:  
          return -np.inf  
      else:  
          return -1.5 * np.log((1 + theta[1]**2) / theta[2])
```

3. log likelihood model

```
[9]: def log_likelihood(theta , x , y , e):  
  
      y_pred = linear_model(x , theta[:2])  
      return -0.5 * np.sum(np.log(2 * np.pi * theta[2]**2) + ((y - y_pred)**2 /  
      ↪theta[2]**2))
```

4. log posterior model

```
[10]: def log_posterior(theta , x , y , e):  
  
      return log_prior(theta) + log_likelihood(theta , x , y , e)
```

```
[11]: n_walkers = 50  
      n_burn = 10000  
      n_steps = 15000  
      n_dim = 3  
  
      starting_guess = np.random.normal(0 , 1 , (n_walkers , n_dim))  
  
      sampler = emcee.EnsembleSampler(n_walkers , n_dim , log_posterior , args = (x ,  
      ↪y , e))  
      sampler.run_mcmc(starting_guess , n_steps)  
  
      chain = sampler.chain[:, n_burn: , :].reshape((-1 , n_dim))
```


/home/coder/.local/lib/python3.10/site-packages/emcee/moves/red_blue.py:99:

RuntimeWarning: invalid value encountered in scalar subtract

```
lnpdiff = f + nlp - state.log_prob[j]
```

```
[12]: plot_mcmc(chain.T[:2 , :] , colors = ['r' , 'r'] , linewidths = [3,3])  
      plt.plot(chain[:,0] , chain[:,1] , 'yo' , alpha = 0.01)
```

```
plt.xlim(-300 , 300)
plt.show()
```



Question 3

Fit the data in Table 1 of arXiv:1008.4686 to a straight line, after including all the data points, (after ignoring x and xy) using both maximum likelihood analysis and using a Bayesian analysis to identify the outliers, using the same procedure as in the second of Jake VanDerPlas blog article. Show graphically the best fit line using both maximum likelihood analysis and also using Bayesian analysis, including the outlier points.

Answer

As we learn in machine learning course about maximum likelihood and bayesian model:

```
[13]: def linear_model(x , theta):

    if len(theta) != 2:
        print("Incorrect size of parameters")
        return

    return theta[0] + theta[1] * x
```

mean squared error (negative log likelihood) for a model

```
[14]: def MSE(theta , x , y , e):

    y_pred = linear_model(x , theta)
    return np.sum((y - y_pred)**2 / (2 * e**2))
```

```
[15]: data = [[201,592,61,],
[244,401,25,],
[47,583,38,],
[287,402,15,],
[203,495,21,],
[58,173,15,],
[210,479,27,],
[202,504,14,],
[198,510,30,],
[158,416,16,],
[165,393,14,],
[201,442,25,],
[157,317,52,],
[131,311,16,],
[166,400,34,],
[160,337,31,],
[186,423,42,],
[125,334,26,],
[218,533,16,],
[146,344,22]]
```

```
data = np.array(data)
```

```
x = data[:, 0]
```

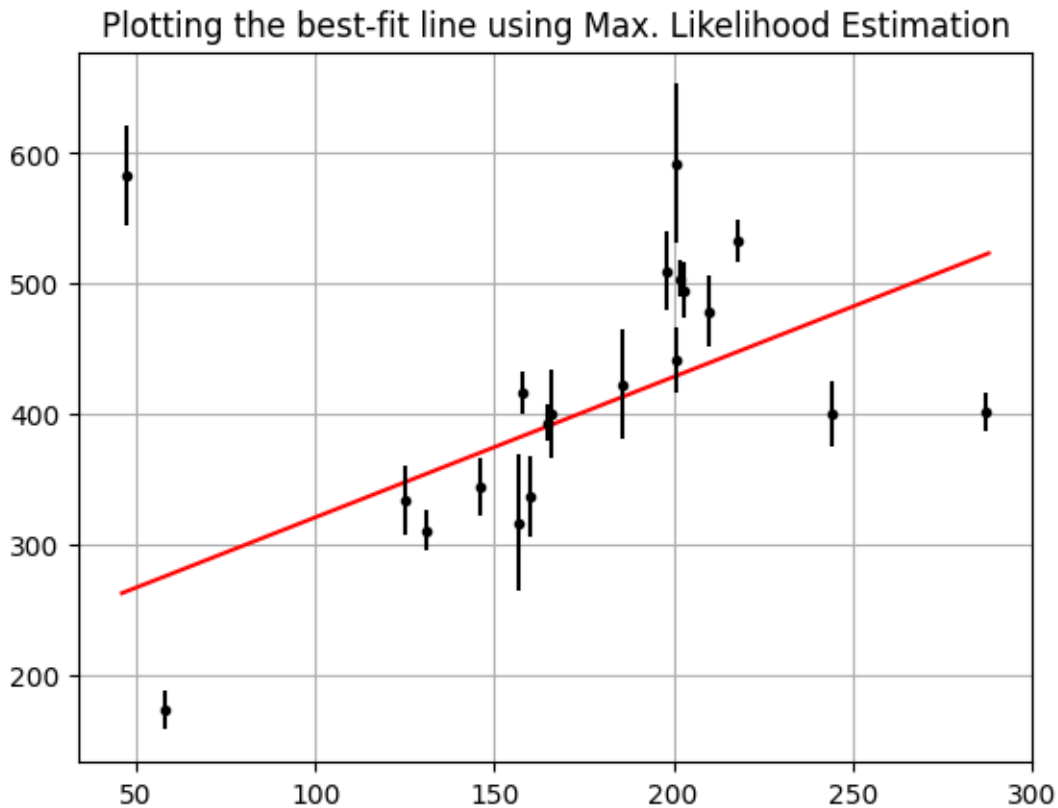
```
y = data[:, 1]
```

```
e = data[:, 2]
```

```
[16]: theta_opt = sp.optimize.fmin(MSE , [0 , 0] , args = (x , y , e) , disp = False)
```

Plotting

```
[17]: x_plot = np.linspace(x.min()-1 , x.max()+1 , 100)
plt.plot(x_plot , linear_model(x_plot , theta_opt) , 'r-' , label = "Linear_
↳Model")
plt.errorbar(x , y , yerr = e , fmt = 'k.')
plt.grid()
plt.title("Plotting the best-fit line using Max. Likelihood Estimation")
plt.show()
```



posterior for a linear model considering Bayesian Analysis

```
[18]: def log_posterior(theta , x , y , e , sigma_B):

    g = theta[2:]
    g = np.clip(g , 0 , 1)

    y_pred = linear_model(x , theta[:2])

    term1 = np.exp(-(y_pred - y)**2 / 2 * e**2) * g / np.sqrt(2 * np.pi * e**2)

    term2 = np.exp(-(y_pred - y)**2 / 2 * sigma_B**2) * (1 - g) / np.sqrt(2 *
↳np.pi * sigma_B**2)
```

```
return np.log(np.product(term1 + term2))
```

```
[19]: n_walkers = 50
      n_steps = 25000
      n_burn = 15000

      theta_guess = np.random.normal(theta_opt , 1 , (n_walkers , 2))
      g_guess = np.random.normal(0.5 , 0.1 , (n_walkers , len(x)))
      starting_guess = np.hstack((theta_guess , g_guess))

      sampler = emcee.EnsembleSampler(nwalkers = n_walkers , ndim = 2+len(x) ,
      ↪ log_prob_fn = log_posterior , args = (x , y , e , 50))
      samples = sampler.run_mcmc(starting_guess , n_steps)

/tmp/ipykernel_3723/3197384969.py:11: RuntimeWarning: divide by zero encountered
in log
    return np.log(np.product(term1 + term2))

[20]: chain = sampler.chain    # shape will be n_walkers x n_steps x n_dim
      final = chain[:, n_burn : , :].reshape(-1 , 2 + len(x)) # we shall remove the
      ↪ n_burn steps

[21]: # calculating the intercept and bias
      theta_bayesian = np.mean(final[:, :2] , axis = 0)
      g = np.mean(final[:, 2:] , axis = 0)
```

Plotting

```
[22]: x_plot = np.linspace(x.min()-1 , x.max()+1 , 100)
      plt.plot(x_plot , linear_model(x_plot , theta_opt) , 'r-' , label = "LM using
      ↪ MLE" , alpha = 1)
      plt.errorbar(x , y , yerr = e , fmt = 'k.')
      plt.plot(x_plot , linear_model(x_plot , theta_bayesian) , 'g-' , label = "LM
      ↪ using Bayesian Analysis")
      plt.legend()
      plt.grid()
      plt.title("Plotting the best-fit line using Max. Likelihood Estimation")
      plt.show()
```


Plotting the best-fit line using Max. Likelihood Estimation

